
Principles of Distributed Database Systems

M. Tamer Özsu
Patrick Valduriez

Outline

- Introduction
- Distributed and parallel database design
- Distributed data control
- Distributed Transaction Processing
- Data Replication
- Database Integration – Multidatabase Systems
- Parallel Database Systems
- Peer-to-Peer Data Management
- Big Data Processing
- NoSQL, NewSQL and Polystores
- Web Data Management

Outline

■ Data Replication

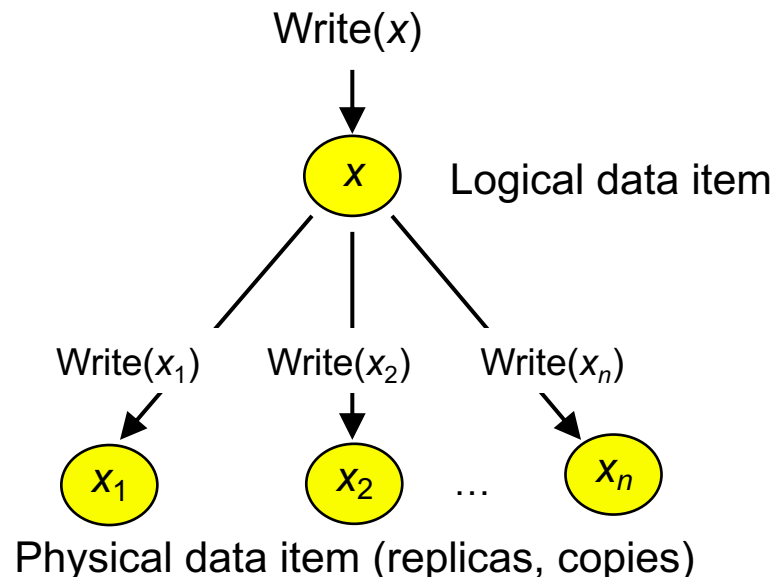
- ❑ Consistency criteria
- ❑ Update Management Strategies
- ❑ Replication Protocols
- ❑ Replication and Failure Management

Replication

- Why replicate?
 - System availability
 - Avoid single points of failure
 - Performance
 - Localization
 - Scalability
 - Scalability in numbers and geographic area
 - Application requirements
- Why not replicate?
 - Replication transparency
 - Consistency issues
 - Updates are costly
 - Availability may suffer if not careful

Execution Model

- There are physical copies of logical objects in the system.
- Operations are specified on logical objects, but translated to operate on physical objects.
- One-copy equivalence
 - Transaction effects on replicated objects should be the same as if they had been performed on a single set of objects.



Replication Issues

- Consistency models - how do we reason about the consistency of the “global execution state”?
 - ❑ Mutual consistency
 - ❑ Transactional consistency
- Where are updates allowed?
 - ❑ Centralized
 - ❑ Distributed
- Update propagation techniques – how do we propagate updates to one copy to the other copies?
 - ❑ Eager
 - ❑ Lazy

Outline

■ Data Replication

- ❑ Consistency criteria
- ❑ Update Management Strategies
- ❑ Replication Protocols
- ❑ Replication and Failure Management

Consistency

■ Mutual Consistency

- ❑ How do we keep the values of physical copies of a logical data item synchronized?
- ❑ Strong consistency
 - All copies are updated within the context of the update transaction
 - When the update transaction completes, all copies have the same value
 - Typically achieved through 2PC
- ❑ Weak consistency
 - Eventual consistency: the copies are not identical when update transaction completes, but they eventually converge to the same value
 - Many versions possible:
 - ❑ Time-bounds
 - ❑ Value-bounds
 - ❑ Drifts

Transactional Consistency

- How can we guarantee that the global execution history over replicated data is serializable?
- One-copy serializability (1SR)
 - The effect of transactions performed by clients on replicated objects should be the same as if they had been performed *one at-a-time* on a single set of objects.
- Weaker forms are possible
 - Snapshot isolation
 - RC-serializability

Example 1

<u>Site A</u>	<u>Site B</u>	<u>Site C</u>
x	x, y	x, y, z
$T_1:$ $x \leftarrow 20$ Write(x) Commit	$T_2:$ Read(x) $x \leftarrow x+y$ Write(y) Commit	$T_3:$ Read(x) Read(y) $z \leftarrow (x*y)/100$ Write(z) Commit

Consider the three histories:

$$H_A = \{W_1(x_A), C_1\}$$

$$H_B = \{W_1(x_B), C_1, R_2(x_B), W_2(y_B), C_2\}$$

$$H_C = \{W_2(y_C), C_2, R_3(x_C), R_3(y_C), W_3(z_C), C_3, W_1(x_C), C_1\}$$

Global history non-serializable: $H_B: T_1 \rightarrow T_2$, $H_C: T_2 \rightarrow T_3 \rightarrow T_1$

Mutually consistent: Assume $x_A = x_B = x_C = 10$, $y_B = y_C = 15$, $y_C = 7$ to begin; in the end
 $x_A = x_B = x_C = 20$, $y_B = y_C = 35$, $y_C = 3.5$

Example 2

<u>Site A</u>	<u>Site B</u>
x	x
T_1 : Read(x) $x \leftarrow x+5$ Write(x) Commit	T_2 : Read(x) $x \leftarrow x*10$ Write(x) Commit

Consider the two histories:

$$H_A = \{R_1(x_A), W_1(x_A), C_1, R_2(x_A), W_2(x_A), C_2\}$$

$$H_B = \{R_2(x_B), W_2(x_B), C_2, R_1(x_B), W_1(x_B), C_1\}$$

Global history non-serializable: $H_A: T_1 \rightarrow T_2, H_B: T_2 \rightarrow T_1$

Mutually inconsistent: Assume $x_A = x_B = 1$ to begin; in the end $x_A = 15, x_B = 60$

Outline

■ Data Replication

- ❑ Consistency criteria
- ❑ Update Management Strategies
- ❑ Replication Protocols
- ❑ Replication and Failure Management


Update Management Strategies

- Depending on when the updates are propagated
 - Eager
 - Lazy
- Depending on where the updates can take place
 - Centralized
 - Distributed

	Centralized	Distributed
Eager		
Lazy		


Eager Replication

- Changes are propagated within the scope of the transaction making the changes. The ACID properties apply to all copy updates.
 - Synchronous
 - Deferred
- ROWA protocol: Read-one/Write-all




Lazy Replication

- Lazy replication first executes the updating transaction on one copy. After the transaction commits, the changes are propagated to all other copies ([refresh transactions](#))
- While the propagation takes place, the copies are mutually inconsistent.
- The time the copies are mutually inconsistent is an adjustable parameter which is application dependent.





Centralized

- There is only one copy which can be updated (the **master**), all others (**slave copies**) are updated reflecting the changes to the master.



Distributed

- Changes can be initiated at any of the copies. That is, any of the sites which owns a copy can update the value of the data item.



Forms of Replication

Eager

- + No inconsistencies (identical copies)
- + Reading the local copy yields the most up to date value
- + Changes are atomic
- A transaction has to update all sites
 - Longer execution time
 - Lower availability

Lazy

- + A transaction is always local (good response time)
- Data inconsistencies
- A local read does not always return the most up-to-date value
- Changes to all copies are not guaranteed
- Replication is not transparent

Centralized

- + No inter-site synchronization is necessary (it takes place at the master)
- + There is always one site which has all the updates
- The load at the master can be high
- Reading the local copy may not yield the most up-to-date value

Distributed

- + Any site can run a transaction
- + Load is evenly distributed
- Copies need to be synchronized

Outline

■ Data Replication

- ❑ Consistency criteria
- ❑ Update Management Strategies
- ❑ Replication Protocols
- ❑ Replication and Failure Management

Replication Protocols

The previous ideas can be combined into 4 different replication protocols:

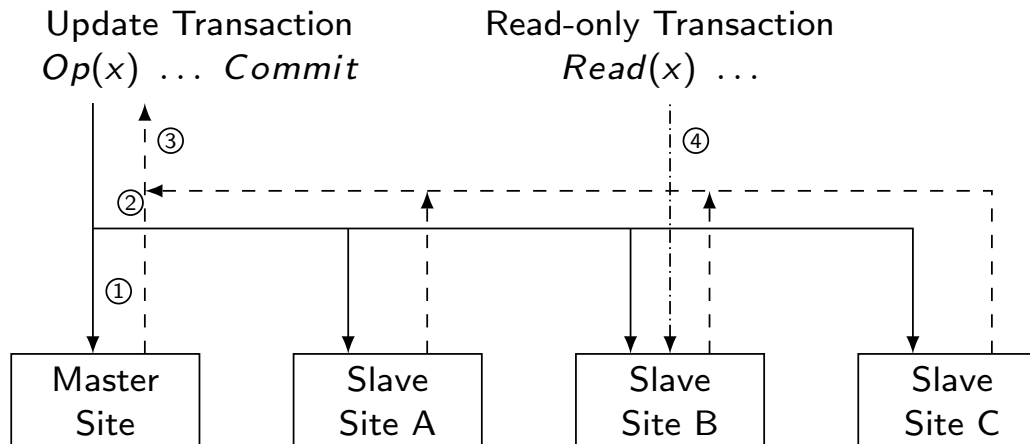
Eager	Eager centralized	Eager distributed
Lazy	Lazy centralized	Lazy distributed
	Centralized	Distributed

Eager Centralized Protocols

- Design parameters:
 - Distribution of master
 - Single master: one master for all data items
 - Primary copy: different masters for different (sets of) data items
 - Level of transparency
 - Limited: applications and users need to know who the master is
 - Update transactions are submitted directly to the master
 - Reads can occur on slaves
 - Full: applications and users can submit anywhere, and the operations will be forwarded to the master
 - Operation-based forwarding
- Four alternative implementation architectures, only three are meaningful:
 - Single master, limited transparency
 - Single master, full transparency
 - Primary copy, full transparency

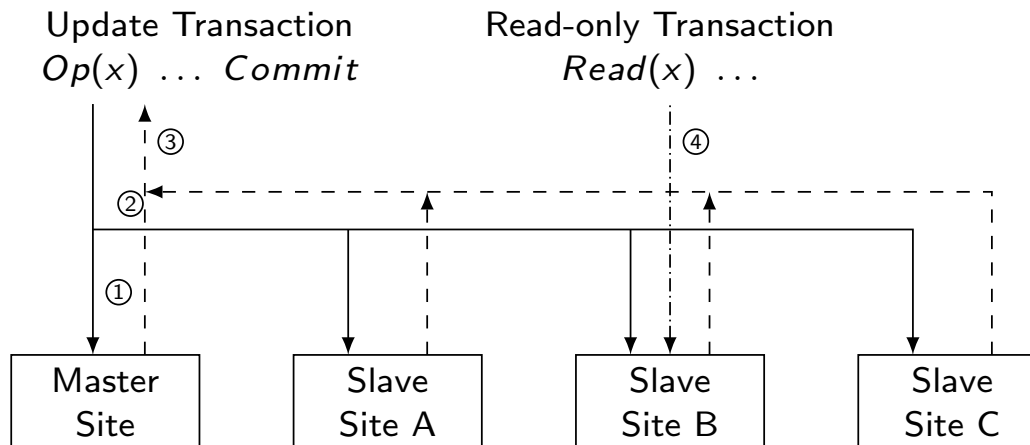
Eager Single Master/Limited Transparency

- Applications submit **update transactions** directly to the master
- Master:
 - Upon read: read locally and return to user
 - Upon write: write locally, multicast write to other replicas (in FFO timestamps order)
 - Upon commit request: run 2PC coordinator to ensure that all have really installed the changes
 - Upon abort: abort and inform other sites about abort
- Slaves install writes that arrive from the master



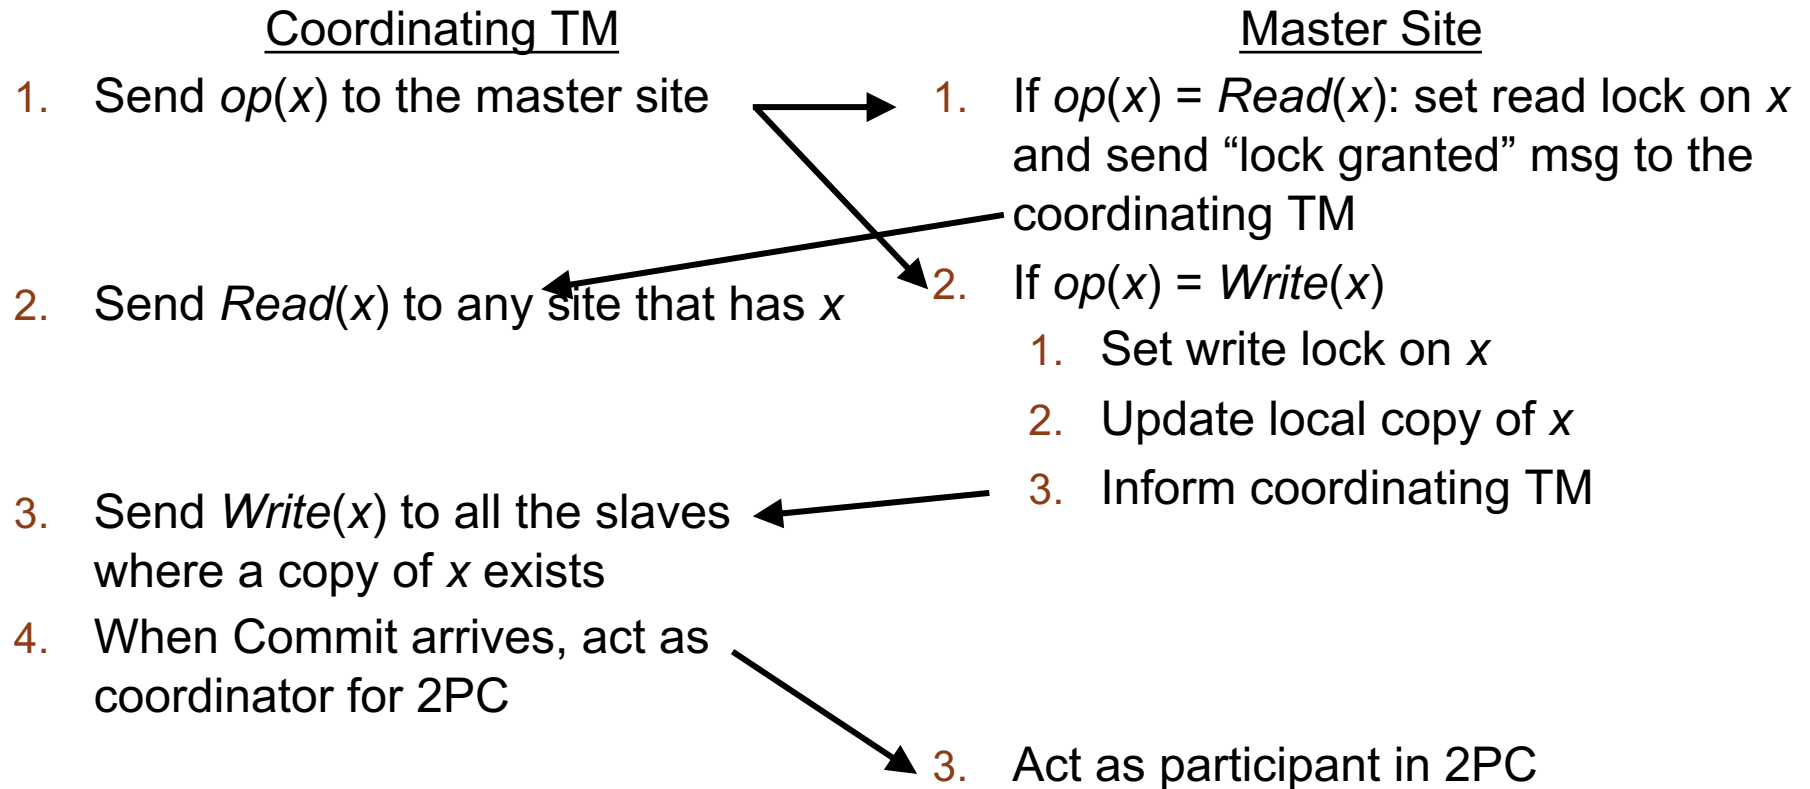
Eager Single Master/Limited Transparency (cont'd)

- Applications submit **read transactions** directly to an appropriate slave
- Slave
 - ❑ Upon read: read locally
 - ❑ Upon write from master copy: execute conflicting writes in the proper order (FIFO or timestamp)
 - ❑ Upon write from client: refuse (abort transaction; there is error)
 - ❑ Upon commit request from read-only: commit locally
 - ❑ Participant of 2PC for update transaction running on primary




Eager Single Master/Full Transparency

Applications submit all transactions to the Transaction Manager at their own sites (Coordinating TM)



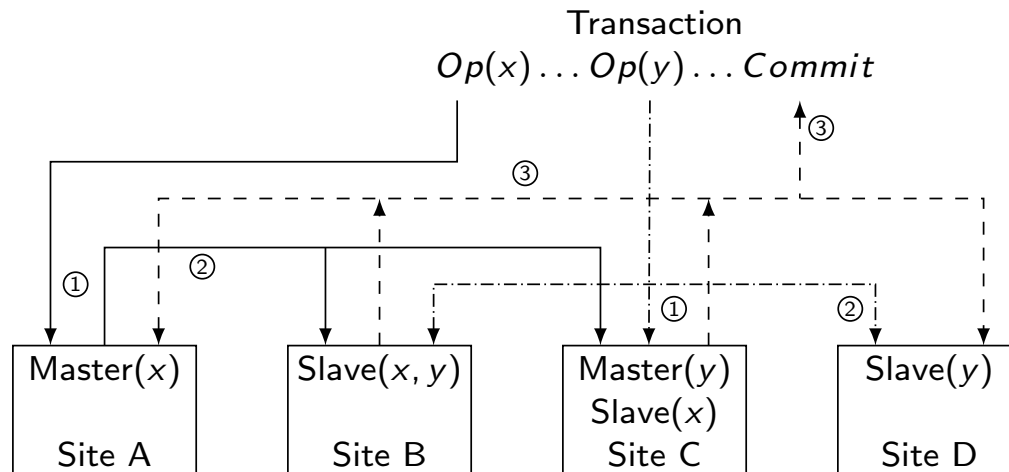
Eager Primary Copy/Full Transparency

- Applications submit transactions directly to their local TMs
- Local TM:
 - ❑ Forward each operation to the primary copy of the data item
 - ❑ Upon granting of locks, submit Read to any slave, Write to all slaves
 - ❑ Coordinate 2PC




Eager Primary Copy/Full Transparency (cont'd)

- Primary copy site
 - ❑ Read(x): lock x and reply to TM
 - ❑ Write(x): lock x , perform update, inform TM
 - ❑ Participate in 2PC
- Slaves: as before



Eager Distributed Protocol

- Updates originate at any copy
 - ❑ Each sites uses 2 phase locking.
 - ❑ Read operations are performed locally.
 - ❑ Write operations are performed at all sites (using a distributed locking protocol).
 - ❑ Coordinate 2PC
- Slaves:
 - ❑ As before




Eager Distributed Protocol (cont'd)

- Critical issue:
 - Concurrent Writes initiated at different master sites are executed in the same order at each slave site
 - Local histories are serializable (this is easy)
- Advantages
 - Simple and easy to implement
- Disadvantage
 - Very high communication overhead
 - n replicas; m update operations in each transaction: $n*m$ messages (assume no multicasting)
 - For throughput of k tps: $k * n*m$ messages
- Alternative
 - Use group communication + deferred update to slaves to reduce messages

Lazy Single Master/Limited Transparency

- Update transactions submitted to master
- Master:
 - Upon read: read locally and return to user
 - Upon write: write locally and return to user
 - Upon commit/abort: terminate locally
 - Sometime after commit: multicast updates to slaves (in order)
- Slaves:
 - Upon read: read locally
 - Refresh transactions: install updates



Lazy Primary Copy/Limited Transparency

- There are multiple masters; each master execution is similar to lazy single master in the way it handles transactions
- Slave execution complicated: refresh transactions from multiple masters and need to be ordered properly

Lazy Primary Copy/Limited Transparency – Slaves

- Assign system-wide unique timestamps to refresh transactions and execute them in timestamp order
 - May cause too many aborts
- Replication graph
 - Similar to serialization graph, but nodes are transactions (T) + sites (S); edge $\langle T_i, S_j \rangle$ exists iff T_i performs a Write(x) and x is stored in S_j
 - For each operation (op_k), enter the appropriate nodes (T_k) and edges; if graph has no cycles, no problem
 - If cycle exists and the transactions in the cycle have been committed at their masters, but their refresh transactions have not yet committed at slaves, abort T_k ; if they have not yet committed at their masters, T_k waits.
- Use group communication

Lazy Single Master/Full Transparency

- This is very tricky
 - Forwarding operations to a master and then getting refresh transactions cause difficulties
- Two problems:
 - Violation of 1SR behavior
 - A transaction may not see its own reads
- Problem arises in primary copy/full transparency as well

Example 3

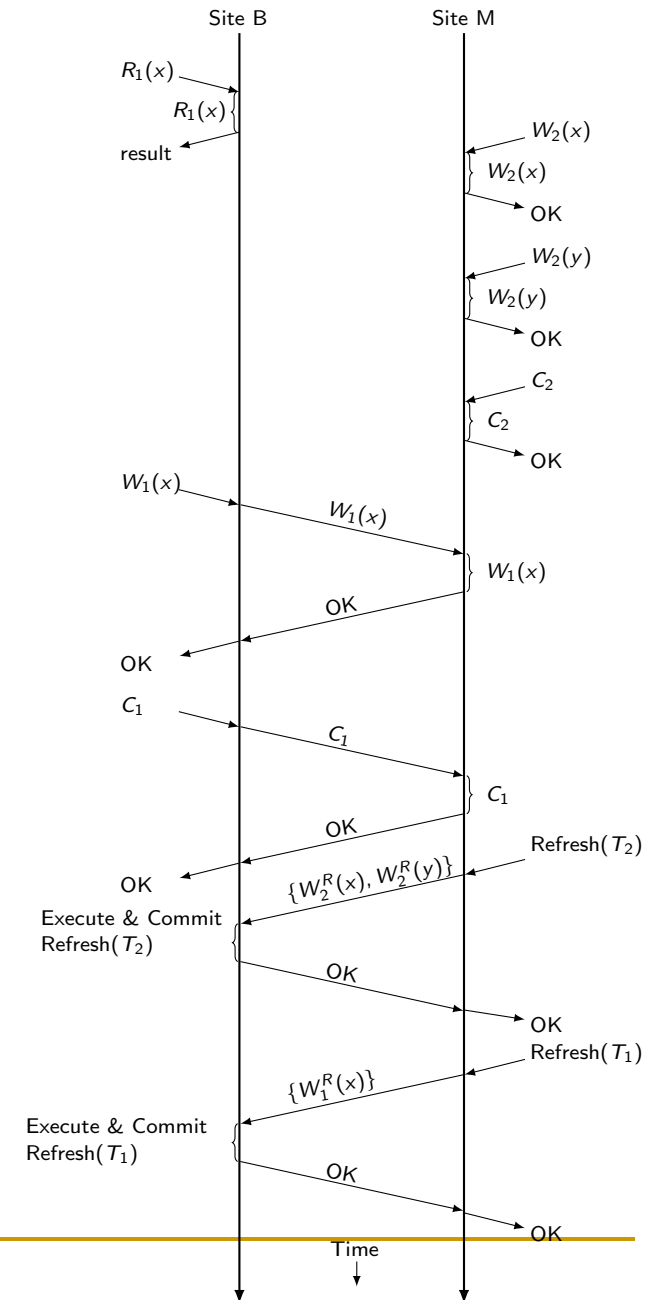
Site M (Master) holds x, y; SiteB holds slave copies of x, y

T_1 : Read(x), Write(y), Commit

T_2 : Read(x), Write(y), Commit

$$H_M = \{W_2(x_M), W_2(y_M), C_2, W_1(y_M), C_1\}$$

$$H_B = \{R_1(x_B), C_1, W_2^R(x_B), W_2^R(y_B), C_2^R, W_1^R(x_B), C_1^R\}$$



Example 4


- Master site M holds x , site C holds slave copy of x
- T_3 : Write(x), Read(x), Commit
- Sequence of execution
 1. $W_3(x)$ submitted at C , forwarded to M for execution
 2. $W_3(x)$ is executed at M , confirmation sent back to C
 3. $R_3(x)$ submitted at C and executed on the local copy
 4. T_3 submits Commit at C , forwarded to M for execution
 5. M executes Commit, sends notification to C , which also commits T_3
 6. M sends refresh transaction for T_3 to C (for $W_3(x)$ operation)
 7. C executes the refresh transaction and commits it
- When C reads x at step 3, it does not see the effects of Write at step 2

Lazy Single Master/ Full Transparency - Solution

- Assume $T = \text{Write}(x)$
- At commit time of transaction T , the master generates a timestamp for it $[ts(T)]$
- Master sets $last_modified(x_M) \leftarrow ts(T)$
- When a refresh transaction arrives at a slave site i , it also sets $last_modified(x_i) \leftarrow last_modified(x_M)$
- Timestamp generation rule at the master:
 - $ts(T)$ should be greater than all previously issued timestamps and should be less than the $last_modified$ timestamps of the data items it has accessed. If such a timestamp cannot be generated, then T is aborted.

Lazy Distributed Replication

- Any site:
 - Upon read: read locally and return to user
 - Upon write: write locally and return to user
 - Upon commit/abort: terminate locally
 - Sometime after commit: send refresh transaction
 - Upon message from other site
 - Detect conflicts
 - Install changes
 - Reconciliation may be necessary



Reconciliation

- Such problems can be solved using pre-arranged patterns:
 - ❑ Latest update win (newer updates preferred over old ones)
 - ❑ Site priority (preference to updates from headquarters)
 - ❑ Largest value (the larger transaction is preferred)
- Or using ad-hoc decision making procedures:
 - ❑ Identify the changes and try to combine them
 - ❑ Analyze the transactions and eliminate the non-important ones
 - ❑ Implement your own priority schemas

Replication Strategies

Eager	<ul style="list-style-type: none">+ Updates do not need to be coordinated+ No inconsistencies- Longest response time- Only useful with few updates- Local copies are can only be read	<ul style="list-style-type: none">+ No inconsistencies+ Elegant (symmetrical solution)- Long response times- Updates need to be coordinated
	<ul style="list-style-type: none">+ No coordination necessary+ Short response times- Local copies are not up to date- Inconsistencies	<ul style="list-style-type: none">+ No centralized coordination+ Shortest response times- Inconsistencies- Updates can be lost (reconciliation)
Centralized		Distributed

Group Communication

- A node can multicast a message to all nodes of a group with a delivery guarantee
- Multicast primitives
 - There are a number of them
 - Total ordered multicast: all messages sent by different nodes are delivered in the same total order at all the nodes
- Used with deferred writes, can reduce communication overhead
 - Remember eager distributed requires $k*m$ messages (with multicast) for throughput of $ktps$ when there are n replicas and m update operations in each transaction
 - With group communication and deferred writes: $2k$ messages

Outline

■ Data Replication

- ❑ Consistency criteria
- ❑ Update Management Strategies
- ❑ Replication Protocols
- Replication and Failure Management

Failures

- So far we have considered replication protocols in the absence of failures
- How to keep replica consistency when failures occur
 - Site failures
 - Read One Write All Available (ROWAA)
 - Communication failures
 - Quorums
 - Network partitioning
 - Quorums

ROWAA with Primary Site

- READ = read any copy, if time-out, read another copy.
- WRITE = send $W(x)$ to **all** copies. If one site rejects the operation, then abort. Otherwise, all sites not responding are “missing writes”.
- VALIDATION = To commit a transaction
 - ❑ Check that all sites in “missing writes” are still down. If not, then abort the transaction.
 - There might be a site recovering concurrent with transaction updates and these may be lost
 - ❑ Check that all sites that were available are still available. If some do not respond, then abort.

Distributed ROWAA

- Each site has a copy of V
 - V represents the set of sites a site believes is available
 - $V(A)$ is the “view” a site has of the system configuration.
- The view of a transaction T [$V(T)$] is the view of its coordinating site, when the transaction starts.
 - Read any copy within V ; update all copies in V
 - If at the end of the transaction the view has changed, the transaction is aborted
- All sites must have the same view!
- To modify V , run a special atomic transaction at all sites.
 - Take care that there are no concurrent views!
 - Similar to commit protocol.
 - Idea: V s have version numbers; only accept new view if its version number is higher than your current one
- Recovery: get missed updates from any active node
 - Problem: no unique sequence of transactions

Quorum-Based Protocol

- Assign a vote to each copy of a replicated object (say V_i) such that $\sum_i V_i = V$
- Each operation has to obtain a read quorum (V_r) to read and a write quorum (V_w) to write an object
- Then the following rules have to be obeyed in determining the quorums:
 - $V_r + V_w > V$ an object is not read and written by two transactions concurrently
 - $V_w > V/2$ two write operations from two transactions cannot occur concurrently on the same object